

Case Study Description

Let us take up the CUSTOMER and TRANSACTIONS table we have created in the Let's Do Together section. Let us solve the following use cases using these tables :-

1. Find out the number of transaction done by each customer (These should be take up in module 8 itself)
2. Create a new table called TRANSACTIONS_COUNT. This table should have 3 fields - custid, fname and count. (Again to be done in module 8)
3. Now write a hive query in such a way that the query populates the data obtained in Step 1 above and populate the table in step 2 above. (This has to be done in module 9).
4. Now lets make the TRANSACTIONS_COUNT table Hbase compliant. In the sence, use Ser Des And Storate handler features of hive to change the TRANSACTIONS_COUNT table to be able to create a TRANSACTIONS table in Hbase. (This has to be done in module 10)
5. Now insert the data in TRANSACTIONS_COUNT table using the query in step 3 again, this should populate the Hbase TRANSACTIONS table automatically (This has to be done in module 10)
6. Now from the Hbase level, write the Hbase java API code to access and scan the TRANSACTIONS table data from java level.

```
h]2.properties Async: true
hive-on-mr is deprecated in Hive 2 and may not be available in the future versions. Consider using a different execution engine (i.e. spark, tez) or using Hive 1.X releases.
hive> show databases;
OK
default
simplidb
test
Time taken: 9.795 seconds, Fetched: 3 row(s)
hive> show tables;
OK
college
customer
olympic
student
Time taken: 0.098 seconds, Fetched: 4 row(s)
hive> use simplidb;
OK
Time taken: 0.042 seconds
hive> CREATE TABLE CUST(Id INT, Firstname STRING, LASTNAME STRING, AGE INT, JOB STRING) row
> format delimited fields terminated by ',';
OK
Time taken: 1.221 seconds
```

```
FAILED: SemanticException [Error 10001]: Table not found: custs
hive> describe cust;
OK
id                int
firstname         string
lastname          string
age               int
job               string
Time taken: 0.156 seconds, Fetched: 5 row(s)
```

```

file:/home/acadgild/Desktop/Assignment_to_be%20submitted/custs.txt
hive> LOAD DATA LOCAL INPATH '/home/acadgild/Desktop/custs.txt' into table CUST;
Loading data to table simplidb.cust
OK
Time taken: 2.724 seconds
hive> select * from cust;
OK
4000001 Kristina Chung 55 Pilot
4000002 Paige Chen 74 Teacher
4000003 Sherri Melton 34 Firefighter
4000004 Gretchen Hill 66 Computer hardware engineer
4000005 Karen Puckett 74 Lawyer
4000006 Patrick Song 42 Veterinarian
4000007 Elsie Hamilton 43 Pilot
4000008 Hazel Bender 63 Carpenter
4000009 Malcolm Wagner 39 Artist
4000010 Dolores McLaughlin 60 Writer
Time taken: 3.988 seconds, Fetched: 10 row(s)

```

```

Hive-on-MR is deprecated in Hive 2 and may not be available in the future versions. Consider using a different execution engine (i.e. spark, tez) or using Hive 1.X releases.
hive> CREATE TABLE TRANSACTIONS(FIRCOL INT, SECCOL STRING,THIRCOL INT,FOU_COL INT,FIF_COL STRING,SIX_COL STRING,SEV_COL STRING,EIG_COL STRING,NIN_COL STRING) row
> format delimited fields terminated by ',';
OK
Time taken: 11.41 seconds
hive> describe TRANSACTIONS;
OK
fircol          int
seccol          string
thircol         int
fou_col         int
fif_col         string
six_col         string
sev_col         string
eig_col         string
nin_col         string
Time taken: 0.533 seconds, Fetched: 9 row(s)
hive>

```

```

Time taken: 0.533 seconds, Fetched: 9 row(s)
hive> LOAD DATA LOCAL INPATH '/home/acadgild/Desktop/txns.txt' into table TRANSACTIONS;
Loading data to table default.transactions
OK
Time taken: 2.488 seconds
hive> select * from TRANSACTIONS;
OK
0 06-26-2011 4000001 40 Exercise & Fitness Cardio Machine Accessories Clarksville Tennessee credit
1 05-26-2011 4000002 198 Exercise & Fitness Weightlifting Gloves Long Beach California credit
2 06-01-2011 4000002 5 Exercise & Fitness Weightlifting Machine Accessories Anaheim California credit
3 06-05-2011 4000003 198 Gymnastics Gymnastics Rings Milwaukee Wisconsin credit
4 12-17-2011 4000002 98 Team Sports Field Hockey Nashville Tennessee credit
5 02-14-2011 4000004 193 Outdoor Recreation Camping & Backpacking & Hiking Chicago Illinois credit
6 10-28-2011 4000005 27 Puzzles Jigsaw Puzzles Charleston South Carolina credit
7 07-14-2011 4000006 96 Outdoor Play Equipment Sandboxes Columbus Ohio credit
8 01-17-2011 4000006 10 Winter Sports Snowmobiling Des Moines Iowa credit
9 05-17-2011 4000006 152 Jumping Bungee Jumping St. Petersburg Florida credit
10 05-29-2011 4000007 180 Outdoor Recreation Archery Reno Nevada credit
11 06-18-2011 4000009 121 Outdoor Play Equipment Swing Sets Columbus Ohio credit
12 02-08-2011 4000009 41 Indoor Games Bowling San Francisco California credit
13 03-13-2011 4000010 107 Team Sports Field Hockey Honolulu Hawaii credit
14 02-25-2011 4000010 36 Gymnastics Vaulting Horses Los Angeles California credit
15 10-20-2011 4000001 137 Combat Sports Fencing Honolulu Hawaii credit
16 05-28-2011 4000010 35 Exercise & Fitness Free Weight Bars Columbia South Carolina credit
17 10-18-2011 4000008 75 Water Sports Scuba Diving & Snorkeling Omaha Nebraska credit
18 11-18-2011 4000008 88 Team Sports Baseball Salt Lake City Utah credit
19 08-29-2011 4000008 51 Water Sports Life Jackets Newark New Jersey credit
20 06-29-2011 4000005 41 Exercise & Fitness Weightlifting Belts New Orleans Louisiana credit
21 02-14-2011 4000005 45 Air Sports Parachutes New York New York credit
22 10-10-2011 4000009 19 Water Sports Kitesurfing Saint Paul Minnesota credit
23 05-02-2011 4000009 99 Gymnastics Gymnastics Rings Springfield Illinois credit
24 06-10-2011 4000003 151 Water Sports Surfing Plano Texas credit
25 10-14-2011 4000009 144 Indoor Games Darts Phoenix Arizona credit
26 10-11-2011 4000009 31 Combat Sports Wrestling Orange California credit
27 09-29-2011 4000010 66 Games Mahjong Fremont California credit
28 05-12-2011 4000008 79 Team Sports Cricket Lexington Kentucky credit
29 06-03-2011 4000001 126 Outdoor Recreation Hunting Phoenix Arizona credit

```

23	05-02-2011	4000009 99	Gymnastics	Gymnastics Rings	Springfield	Illinois	credit	
24	06-10-2011	4000003 151	Water Sports	Surfing Plano	Texas		credit	
25	10-14-2011	4000009 144	Indoor Games	Darts Phoenix	Arizona		credit	
26	10-11-2011	4000009 31	Combat Sports	Wrestling Orange	California		credit	
27	09-29-2011	4000010 66	Games	Mahjong Fremont	California		credit	
28	05-12-2011	4000008 79	Team Sports	Cricket Lexington	Kentucky		credit	
29	06-03-2011	4000001 126	Outdoor Recreation	Hunting Phoenix	Arizona		credit	
30	03-14-2011	4000001 47	Water Sports	Swimming Lincoln	Nebraska		credit	
31	11-28-2011	4000008 5	Games	Dice & Dice Sets Los Angeles	California		credit	
32	01-29-2011	4000008 20	Team Sports	Soccer Springfield	Illinois		credit	
33	06-15-2011	4000008 154	Outdoor Recreation	Lawn Games Nashville	Tennessee		credit	
34	05-06-2011	4000008 98	Team Sports	Indoor Volleyball Atlanta	Georgia		credit	
35	04-12-2011	4000008 185	Games	Board Games Centennial	Colorado		credit	
36	10-13-2011	4000007 35	Team Sports	Football Saint Paul	Minnesota		credit	
37	04-19-2011	4000007 20	Outdoor Recreation	Shooting Games San Diego	California		credit	
38	08-05-2011	4000007 150	Outdoor Recreation	Camping & Backpacking & Hiking Hampton	Virginia		credit	
39	03-12-2011	4000006 174	Outdoor Play Equipment	Swing Sets Pittsburgh	Pennsylvania		credit	
40	11-07-2011	4000005 165	Team Sports	Cheerleading Reno	Nevada		credit	
41	04-16-2011	4000004 28	Indoor Games	Bowling Westminster	Colorado		cash	
42	09-10-2011	4000004 38	Outdoor Recreation	Tetherball Denton	Texas		cash	
43	04-22-2011	4000004 32	Water Sports	Water Polo Las Vegas	Nevada		cash	
44	09-11-2011	4000001 135	Water Sports	Surfing Seattle	Washington		credit	
45	11-27-2011	4000001 90	Exercise & Fitness	Abdominal Equipment Honolulu	Hawaii		credit	
46	05-27-2011	4000001 52	Gymnastics	Vaulting Horses Cleveland	Ohio		credit	
47	10-23-2011	4000008 100	Outdoor Play Equipment	Swing Sets Everett	Washington		credit	
48	09-27-2011	4000007 157	Exercise & Fitness	Exercise Bands Philadelphia	Pennsylvania		credit	
49	07-12-2011	4000010 144	Jumping	Jumping Stilts Cambridge	Massachusetts		credit	
50	10-20-2011	4000010 55	Jumping	Pogo Sticks Everett	Washington		credit	
51	02-17-2011	4000002 32	Water Sports	Life Jackets Columbus	Georgia		cash	
52	02-04-2011	4000005 44	Outdoor Play Equipment	Lawn Water Slides Hampton	Virginia		cash	
53	06-12-2011	4000004 44	Water Sports	Scuba Diving & Snorkeling Charleston	South Carolina		cash	
54	10-03-2011	4000007 154	Outdoor Recreation	Running Long Beach	California		credit	
55	12-16-2011	4000006 106	Water Sports	Swimming New York	New York		credit	
56	06-21-2011	4000002 176	Outdoor Recreation	Geocaching Boston	Massachusetts		credit	
57	12-20-2011	4000003 178	Outdoor Recreation	Skating San Jose	California		credit	
58	12-29-2011	4000002 194	Water Sports	Windsurfing Oklahoma City	Oklahoma		credit	
59	11-07-2011	4000001 21	Winter Sports	Snowboarding Philadelphia	Pennsylvania		cash	

Time taken: 3.878 seconds, Fetched: 60 row(s)

rhaxTerm hv e shorritin to the pnfessinal editon here: <https://lnhaxterm.mnhatak.net>

1. Find out the number of transaction done by each customer (These should be take up in module 8 itself)

Solution:

**select a.custid, a.fname, count(a.fname) from CUSTOMER a join
TXNRECORDS b on a.custid =b.custno group by a.fname,a.custid;**

```
4000010 Doreen McLaughlin 00 Writer
Time taken: 0.876 seconds, Fetched: 10 row(s)
hive> select a.id, a.Firstname, count(a.Firstname) from CUST a join TRANSACTIONS b on a.id =b.THIRCOL group by a.Firstname,a.id;
WARNING: Hive-on-MR is deprecated in Hive 2 and may not be available in the future versions. Consider using a different execution engine
(i.e. spark, tez) or using Hive 1.X releases.
Query ID = acadgild_20180926114911_0b240829-c845-40cc-9271-804ff7831f5d
Total jobs = 1
SLF4J: Class path contains multiple SLF4J bindings.
SLF4J: Found binding in [jar:file:/home/acadgild/install/hive/apache-hive-2.3.2-bin/lib/log4j-slf4j-impl-2.6.2.jar!/org/slf4j/impl/Static
LoggerBinder.class]
SLF4J: Found binding in [jar:file:/home/acadgild/install/hadoop/hadoop-2.6.5/share/hadoop/common/lib/slf4j-log4j12-1.7.5.jar!/org/slf4j/i
mpl/StaticLoggerBinder.class]
SLF4J: See http://www.slf4j.org/codes.html#multiple_bindings for an explanation.
SLF4J: Actual binding is of type [org.apache.logging.slf4j.Log4jLoggerFactory]
2018-09-26 11:49:28 Starting to launch local task to process map join; maximum memory = 518979584
2018-09-26 11:49:32 Dump the side-table for tag: 0 with group count: 10 into file: file:/tmp/acadgild/93e76ba3-a7f0-428f-a752-b9d8396
199a2/hive_2018-09-26_11-49-11_175_5355354417705488424-1/-local-10005/HashTable-Stage-2/MapJoin-mapfile00--.hashtable
2018-09-26 11:49:32 Uploaded 1 File to: file:/tmp/acadgild/93e76ba3-a7f0-428f-a752-b9d8396199a2/hive_2018-09-26_11-49-11_175_53553544
17705488424-1/-local-10005/HashTable-Stage-2/MapJoin-mapfile00--.hashtable (556 bytes)
2018-09-26 11:49:32 End of local task; Time Taken: 3.901 sec.
Execution completed successfully
MapredLocal task succeeded
Launching Job 1 out of 1
Number of reduce tasks not specified. Estimated from input data size: 1
In order to change the average load for a reducer (in bytes):
  set hive.exec.reducers.bytes.per.reducer=<number>
In order to limit the maximum number of reducers:
  set hive.exec.reducers.max=<number>
In order to set a constant number of reducers:
  set mapreduce.job.reduces=<number>
Starting Job = job_1537939271077_0001, Tracking URL = http://localhost:8088/proxy/application_1537939271077_0001/
Kill Command = /home/acadgild/install/hadoop/hadoop-2.6.5/bin/hadoop job -kill job_1537939271077_0001
Hadoop job information for Stage-2: number of mappers: 1; number of reducers: 1
2018-09-26 11:49:58,526 Stage-2 map = 0%, reduce = 0%
2018-09-26 11:50:16,978 Stage-2 map = 100%, reduce = 0%, Cumulative CPU 4.48 sec
2018-09-26 11:50:34,329 Stage-2 map = 100%, reduce = 100%, Cumulative CPU 7.3 sec
```

```

hive command: /home/abogato/2nd/etl/hadoop/hadoop-2.6.0/bin/hadoop job -kill job_1537939271077_0001
Hadoop job information for Stage-2: number of mappers: 1; number of reducers: 1
2018-09-26 11:49:58.526 Stage-2 map = 0%, reduce = 0%
2018-09-26 11:50:16.978 Stage-2 map = 100%, reduce = 0%, Cumulative CPU 4.48 sec
2018-09-26 11:50:34.329 Stage-2 map = 100%, reduce = 100%, Cumulative CPU 7.3 sec
MapReduce Total cumulative CPU time: 7 seconds 300 msec
Ended Job = job_1537939271077_0001
MapReduce Jobs Launched:
Stage-Stage-2: Map: 1 Reduce: 1 Cumulative CPU: 7.3 sec HDFS Read: 17843 HDFS Write: 381 SUCCESS
Total MapReduce CPU Time Spent: 7 seconds 300 msec
OK
4000001 Kristina 8
4000002 Paige 6
4000003 Sherri 3
4000004 Gretchen 5
4000005 Karen 5
4000006 Patrick 5
4000007 Elsie 6
4000008 Hazel 10
4000009 Malcolm 6
4000010 Dolores 6
Time taken: 84.397 seconds, Fetched: 10 row(s)
hive>

```

2. Create a new table called **TRANSACTIONS_COUNT**. This table should have 3 fields - custid, fname and count. (Again to be done in module 8)

Solution:

create table **TRANSACTIONS_COUNT**(custid int, fname String, count string);

```

4000008 Hazel 10
4000009 Malcolm 6
4000010 Dolores 6
Time taken: 84.397 seconds, Fetched: 10 row(s)
hive> create table TRANSACTIONS_COUNT(custid int, fname String, count string);
OK
Time taken: 0.481 seconds
hive>

```

3. Now write a hive query in such a way that the query populates the data obtained in Step 1 above and populate the table in step 2 above. (This has to be done in module 9).

Solution:

INSERT OVERWRITE TABLE TRANSACTIONS_COUNT select a.custid, a.fname, count(a.fname) from CUSTOMER a join TXNRECORDS b on a.custid =b.custno group by a.fname,a.custid;

```

FAILED: SemanticException [Error 10002]: Line 1:57 Invalid column reference 'fname'
hive> INSERT OVERWRITE TABLE TRANSACTIONS_COUNT select a.id, a.Firstname, count(a.Firstname) from CUST a join TRANSACTIONS b on a.id =b.T
HRCOL group by a.Firstname,a.id;
WARNING: Hive-on-MR is deprecated in Hive 2 and may not be available in the future versions. Consider using a different execution engine
(i.e. spark, tez) or using Hive 1.X releases.
Query ID = acadgild_20180926115702_6653e533-60be-42cf-8d95-1c3675dd0272
Total jobs = 1
SLF4J: Class path contains multiple SLF4J bindings.
SLF4J: Found binding in [jar:file:/home/acadgild/install/hive/apache-hive-2.3.2-bin/lib/log4j-slf4j-impl-2.6.2.jar!/org/slf4j/impl/Static
LoggerBinder.class]
SLF4J: Found binding in [jar:file:/home/acadgild/install/hadoop/hadoop-2.6.5/share/hadoop/common/lib/slf4j-log4j12-1.7.5.jar!/org/slf4j/i
mpl/StaticLoggerBinder.class]
SLF4J: See http://www.slf4j.org/codes.html#multiple_bindings for an explanation.
SLF4J: Actual binding is of type [org.apache.logging.slf4j.Log4jLoggerFactory]
2018-09-26 11:57:17 Starting to launch local task to process map join; maximum memory = 518979584
2018-09-26 11:57:22 Dump the side-table for tag: 0 with group count: 10 into file: file:/tmp/acadgild/93e76ba3-a7f0-428f-a752-b9d8396
199a2/hive_2018-09-26_11-57-02_671_6414821775807981277-1/-local-10003/HashTable-Stage-2/MapJoin-mapfile10--.hashtable
2018-09-26 11:57:22 Uploaded 1 File to: file:/tmp/acadgild/93e76ba3-a7f0-428f-a752-b9d8396199a2/hive_2018-09-26_11-57-02_671_64148217
75807981277-1/-local-10003/HashTable-Stage-2/MapJoin-mapfile10--.hashtable (556 bytes)
2018-09-26 11:57:22 End of local task; Time Taken: 4.255 sec.
Execution completed successfully
MapredLocal task succeeded
Launching Job 1 out of 1
Number of reduce tasks not specified. Estimated from input data size: 1
In order to change the average load for a reducer (in bytes):
  set hive.exec.reducers.bytes.per.reducer=<number>
In order to limit the maximum number of reducers:
  set hive.exec.reducers.max=<number>
In order to set a constant number of reducers:
  set mapreduce.job.reduces=<number>
Starting Job = job_1537939271077_0002, Tracking URL = http://localhost:8088/proxy/application_1537939271077_0002/
Kill Command = /home/acadgild/install/hadoop/hadoop-2.6.5/bin/hadoop job -kill job_1537939271077_0002
Hadoop job information for Stage-2: number of mappers: 1; number of reducers: 1
2018-09-26 11:57:41,868 Stage-2 map = 0%, reduce = 0%
2018-09-26 11:57:55,085 Stage-2 map = 100%, reduce = 0%, Cumulative CPU 4.06 sec
2018-09-26 11:58:13,438 Stage-2 map = 100%, reduce = 100%, Cumulative CPU 9.05 sec
MapReduce Total cumulative CPU time: 9 seconds 50 msec

```

```

set hive.exec.reducers.max=<number>
In order to set a constant number of reducers:
  set mapreduce.job.reduces=<number>
Starting Job = job_1537939271077_0002, Tracking URL = http://localhost:8088/proxy/application_1537939271077_0002/
Kill Command = /home/acadgild/install/hadoop/hadoop-2.6.5/bin/hadoop job -kill job_1537939271077_0002
Hadoop job information for Stage-2: number of mappers: 1; number of reducers: 1
2018-09-26 11:57:41,868 Stage-2 map = 0%, reduce = 0%
2018-09-26 11:57:55,085 Stage-2 map = 100%, reduce = 0%, Cumulative CPU 4.06 sec
2018-09-26 11:58:13,438 Stage-2 map = 100%, reduce = 100%, Cumulative CPU 9.05 sec
MapReduce Total cumulative CPU time: 9 seconds 50 msec
Ended Job = job_1537939271077_0002
Loading data to table default.transactions_count
MapReduce Jobs Launched:
Stage-Stage-2: Map: 1 Reduce: 1 Cumulative CPU: 9.05 sec HDFS Read: 18492 HDFS Write: 257 SUCCESS
Total MapReduce CPU Time Spent: 9 seconds 50 msec
OK
Time taken: 72.988 seconds
hive> select * from TRANSACTIONS_COUNT;
OK
4000001 Kristina 8
4000002 Paige 6
4000003 Sherri 3
4000004 Gretchen 5
4000005 Karen 5
4000006 Patrick 5
4000007 Elsie 6
4000008 Hazel 10
4000009 Malcolm 6
4000010 Dolores 6
Time taken: 0.35 seconds, Fetched: 10 row(s)
hive>

```

4. Now let's make the TRANSACTIONS_COUNT table Hbase compliant. In the process, use SerDes and Storage handler features of Hive to change the TRANSACTIONS_COUNT table to be able to create a TRANSACTIONS table in Hbase. (This has to be done in module 10)

Solution:

```

create table TRANSACTIONS_COUNT(custid int, fname String, count string)
STORED BY 'org.apache.hadoop.hive.hbase.HBaseStorageHandler' with serdeproperties
("hbase.columns.mapping"=":key,customerdetails:fname, customerdetails:count")
tblproperties("hbase.table.name"="transactions_count");

```

```

mycustomer_ext
transactions
transactions_count
txnrecords
users
Time taken: 0.088 seconds, Fetched: 14 row(s)
hive> create table TRANSACTIONS_COUNT(custid int, fname String, count string) STORED BY 'org.apache.hadoop.hive.hbase.HBaseStorageHandler'
'with serdeproperties ("hbase.columns.mapping"=":key,customerdetails:fname, customerdetails:count") tblproperties("hbase.table.name"="tra
nsactions_count");

```

5. Now insert the data in TRANSACTIONS_COUNT table using the query in step 3 again, this should populate the Hbase TRANSACTIONS table automatically (This has to be done in module 10)

Solution :

```

hive> > INSERT OVERWRITE TABLE TRANSACTIONS_COUNT select a.id, a.Firstname, count(a.Firstname) from CUST a join TRANSACTIONS b on a.id =b.T
HIRCOL group by a.Firstname,a.id;
WARNING: Hive-on-MR is deprecated in Hive 2 and may not be available in the future versions. Consider using a different execution engine
(i.e. spark, tez) or using Hive 1.X releases.
Query ID = acadgild_20180929185348_f4601a61-b0eb-441a-8ec6-3877356a0bab
Total jobs = 1
SLF4J: Class path contains multiple SLF4J bindings.
SLF4J: Found binding in [jar:file:/home/acadgild/install/hive/apache-hive-2.3.2-bin/lib/log4j-slf4j-impl-2.6.2.jar!/org/slf4j/impl/Static
LoggerBinder.class]
SLF4J: Found binding in [jar:file:/home/acadgild/install/hadoop/hadoop-2.6.5/share/hadoop/common/lib/slf4j-log4j12-1.7.5.jar!/org/slf4j/i
mpl/StaticLoggerBinder.class]
SLF4J: See http://www.slf4j.org/codes.html#multiple_bindings for an explanation.
SLF4J: Actual binding is of type [org.apache.logging.slf4j.Log4jLoggerFactory]
2018-09-29 18:54:06 Starting to launch local task to process map join; maximum memory = 518979584
2018-09-29 18:54:10 Dump the side-table for tag: 0 with group count: 10 into file: file:/tmp/acadgild/6fd6525c-6fc7-4db3-b7d8-bddecbc
ca9e0/hive_2018-09-29_18-53-48_107_4923524814062397576-1/-local-10003/HashTable-Stage-2/MapJoin-mapfile000--.hashtable
2018-09-29 18:54:10 Uploaded 1 File to: file:/tmp/acadgild/6fd6525c-6fc7-4db3-b7d8-bddecbc9e0/hive_2018-09-29_18-53-48_107_49235248
14062397576-1/-local-10003/HashTable-Stage-2/MapJoin-mapfile000--.hashtable (556 bytes)
2018-09-29 18:54:10 End of local task; Time Taken: 4.103 sec.
Execution completed successfully
MapredLocal task succeeded
Launching Job 1 out of 1
Number of reduce tasks not specified. Estimated from input data size: 1
In order to change the average load for a reducer (in bytes):
  set hive.exec.reducers.bytes.per.reducer=<number>
In order to limit the maximum number of reducers:
  set hive.exec.reducers.max=<number>
In order to set a constant number of reducers:
  set mapreduce.job.reduces=<number>
Starting Job = job_1538223396573_0003, Tracking URL = http://localhost:8088/proxy/application_1538223396573_0003/
Kill Command = /home/acadgild/install/hadoop/hadoop-2.6.5/bin/hadoop job -kill job_1538223396573_0003
Hadoop job information for Stage-2: number of mappers: 1; number of reducers: 1
2018-09-29 18:54:30,204 Stage-2 map = 0%, reduce = 0%
2018-09-29 18:54:46,510 Stage-2 map = 100%, reduce = 0%, Cumulative CPU 4.64 sec
2018-09-29 18:55:03,837 Stage-2 map = 100%, reduce = 100%, Cumulative CPU 9.53 sec
MapReduce Total cumulative CPU time: 9 seconds 530 msec
Ended Job = job_1538223396573_0003

```

```

Kill Command = /home/acadgild/install/hadoop/hadoop-2.6.5/bin/hadoop job -kill job_1538223396573_0003
Hadoop job information for Stage-2: number of mappers: 1; number of reducers: 1
2018-09-29 18:54:30,204 Stage-2 map = 0%, reduce = 0%
2018-09-29 18:54:46,510 Stage-2 map = 100%, reduce = 0%, Cumulative CPU 4.64 sec
2018-09-29 18:55:03,837 Stage-2 map = 100%, reduce = 100%, Cumulative CPU 9.53 sec
MapReduce Total cumulative CPU time: 9 seconds 530 msec
Ended Job = job_1538223396573_0003
Loading data to table simplidb.transactions_count
MapReduce Jobs Launched:
Stage-Stage-2: Map: 1 Reduce: 1 Cumulative CPU: 9.53 sec HDFS Read: 18619 HDFS Write: 258 SUCCESS
Total MapReduce CPU Time Spent: 9 seconds 530 msec
OK
Time taken: 78.424 seconds
hive> select * from transactions_count;
OK
4000001 Kristina 8
4000002 Paige 6
4000003 Sherri 3
4000004 Gretchen 5
4000005 Karen 5
4000006 Patrick 5
4000007 Elsie 6
4000008 Hazel 10
4000009 Malcolm 6
4000010 Dolores 6
Time taken: 0.515 seconds, Fetched: 10 row(s)
hive>

```


6. Now from the Hbase level, write the Hbase java API code to access and scan the TRANSACTIONS table data from java level

Solution:

```
import org.apache.hadoop.conf.Configuration;
import org.apache.hadoop.hbase.HBaseConfiguration;
import org.apache.hadoop.hbase.client.HTable;
import org.apache.hadoop.hbase.client.Result;
import org.apache.hadoop.hbase.client.ResultScanner;
import org.apache.hadoop.hbase.client.Scan;
import org.apache.hadoop.hbase.util.Bytes;

public class Usecase1 {
    public static void main(String[] args) throws IOException, InterruptedException {
        Configuration conf = HBaseConfiguration.create();

        //System.out.println("Creating HTable instance ");
        HTable table = new HTable(conf, "transactions_count");

        System.out.println("Creating scan object to scan column family customer details");
        //Scan scan = new Scan(Bytes.toBytes("john"), Bytes.toBytes("p4"));
        Scan scan = new Scan();
        scan.addFamily(Bytes.toBytes("customerdetails"));
        System.out.println("Getting a result scanner object...");
        ResultScanner rs = table.getScanner(scan);

        for (Result r : rs) {
            //System.out.println("Result: " + r);
        }

        rs.close();
    }
}
```